

Volumes and Data

Exercise 9.1: Create a ConfigMap

Overview

Container files are ephemeral, which can be problematic for some applications. Should a container be restarted the files will be lost. In addition, we need a method to share files between containers inside a Pod.

A Volume is a directory accessible to containers in a Pod. Cloud providers offer volumes which persist further than the life of the Pod, such that AWS or GCE volumes can be pre-populated and offered to Pods, or transferred from one Pod to another. Config is also another popular solution for dynamic, persistent volumes.

We are here!

Online

Unlike current Docker volumes a Kubernetes volume has the lifetime of the Pod, not the containers within. You can also use different types of volumes in the same Pod simultaneously, but Volumes cannot mount in a nested fashion. Each must have their own mount point. Volumes are declared with `spec.volumes` and mount points with `spec.containers.volumeMounts` parameters. Each particular volume type, 24 currently, may have other restrictions. <https://kubernetes.io/docs/concepts/storage/volumes/#types-of-volumes>

We will also work with a ConfigMap, which is basically a set of key-value pairs. This data can be made available so that a Pod can read the data as environment variables or configuration data. A ConfigMap is similar to a Secret, except they are not base64 byte encoded values. They



There are three different ways a ConfigMap can ingest data, from a literal value, from a file or from a directory of files

1. We will create a ConfigMap containing primary colors. We will create a series of files to ingest into the ConfigMap. First, we create a directory primary and populate it with four files. Then we create a file in our home directory with our favorite color.

```
1
2      student@lfs458-node-1a0a:~$ mkdir primary
3
4      student@lfs458-node-1a0a:~$ echo c > primary/cyan
5
6      student@lfs458-node-1a0a:~$ echo m > primary/magenta
7
8      student@lfs458-node-1a0a:~$ echo y > primary/yellow
9
10     student@lfs458-node-1a0a:~$ echo k > primary/black
11
12     student@lfs458-node-1a0a:~$ echo "known as key" >>
13 primary/black
14
15     student@lfs458-node-1a0a:~$ echo blue > favorite
16
```

2. Now we will create the ConfigMap and populate it with the files we created as well as a literal value from the command line.

```
1
2      student@lfs458-node-1a0a:~$ kubectl create configmap
3 colors \
4      --from-literal=text=black \
5      --from-file=./favorite \
6      --from-file=./primary/
7 configmap/colors created
8
```

3. View how the data is organized inside the cluster.

```

1
2      student@lfs458-node-1a0a:~$ kubectl get configmap
3 colors
4      NAME    DATA    AGE
5      colors  6        30s
6
7      student@lfs458-node-1a0a:~$ kubectl get configmap
8 colors -o yaml
9      apiVersion: v1
10     data:
11         black: |
12             k
13             known as key
14         cyan: |
15             c
16         favorite: |
17             blue
18         magenta: |
19             m
20         text: black
21         yellow: |
22             y
23     kind: ConfigMap
24     <output_omitted>

```

4. Now we can create a Pod to use the ConfigMap. In this case a particular parameter is being defined as an environment variable.

```

1
2      student@lfs458-node-1a0a:~$ vim simpleshell.yaml
3      apiVersion: v1
4      kind: Pod
5      metadata:
6          name: shell-demo
7      spec:
8      containers:
9      - name: nginx
10        image: nginx
11        env:
12        - name: ilike
13            valueFrom:
14                configMapKeyRef:
15                    name: colors
16                    key:
17 favorite
18
19

```

5. Create the Pod and view the environmental variable. After you view the parameter, exit out and delete the pod.

```
1
2      student@lfs458-node-1a0a:~$ kubectl create -f
3  simpleshell.yaml
4      pod/shell-demo created
5
6      student@lfs458-node-1a0a:~$ kubectl exec -it shell-
7  demo \
8      -- /bin/bash -c 'echo $ilike'
9
10     blue
11
12      student@lfs458-node-1a0a:~$ kubectl delete pod
13  shell-demo
      pod "shell-demo" deleted
```

6. All variables from a file can be included as environment variables as well. Comment out the previous env: stanza and add a slightly different envFrom to the file. Having new and old code at the same time can be helpful to see and understand the differences. Recreate the Pod, check all variables and delete the pod again. They can be found spread throughout the environment variable output.

```
1
2      student@lfs458-node-1a0a:~$ vim simpleshell.yaml
3      <output_omitted>
4          image: nginx
5      #      env:
6      #      - name: ilike
7      #      valueFrom:
8      #      configMapKeyRef:
9      #      name: colors
10     #      key: favorite
11     envFrom:
12     - configMapRef:
13     name: colors
14     student@lfs458-node-1a0a:~$ kubectl create -f
15     simpleshell.yaml
16     pod/shell-demo created
17
18     student@lfs458-node-1a0a:~$ kubectl exec -it shell-
19     demo \
20         -- /bin/bash -c 'env'
21     HOSTNAME=shell-demo
22     NJS_VERSION=1.13.6.0.1.14-1~stretch
23     NGINX_VERSION=1.13.6-1~stretch
24     black=k
25     know as key
26
27     favorite=blue
28     <output_omitted>
29
30     student@lfs458-node-1a0a:~$ kubectl delete pod
31     shell-demo
32     pod "shell-demo" deleted
```

7. A ConfigMap can also be created from a YAML file. Create one with a few parameters to describe a car.

```
1
2     student@lfs458-node-1a0a:~$ vim car-map.yaml
3     apiVersion: v1
4     kind: ConfigMap
5     metadata:
6         name: fast-car
7         namespace: default
8     data:
9         car.make: Ford
10        car.model: Mustang
11        car.trim: Shelby
12
13
```

8. Create the ConfigMap and verify the settings.

```
1
2     student@lfs458-node-1a0a:~$ kubectl create -f car-
3 map.yaml
4     configmap/fast-car created
5
6     student@lfs458-node-1a0a:~$ kubectl get configmap
7 fast-car -o yaml
8     apiVersion: v1
9     data:
10        car.make: Ford
11        car.model: Mustang
12        car.trim: Shelby
13     kind: ConfigMap
14     <output_omitted>
```

9. We will now make the ConfigMap available to a Pod as a mounted volume. You can again comment out the previous environmental settings and add the following new stanza. The containers: and volumes: entries are indented the same number of spaces.

```
1
2      student@lfs458-node-1a0a:~$ vim simpleshell.yaml
3      <output_omitted>
4      spec:
5          containers:
6              - name: nginx
7                image: nginx
8                volumeMounts:
9                  - name: car-vol
10                    mountPath: /etc/cars
11      volumes:
12          - name: car-vol
13            configMap:
14              name: fast-car
15      <comment out rest of file>
16
17
18
```

10. Create the Pod again. Verify the volume exists and the contents of a file within. Due to the lack of a carriage return in the file your next prompt may be on the same line as the output, Shelby.

```
1
2      student@lfs458-node-1a0a:~$ kubectl create -f
3      simpleshell.yaml
4      pod "shell-demo" created
5
6      student@lfs458-node-1a0a:~$ kubectl exec -it shell-
7      demo -- \
8          /bin/bash -c 'df -ha |grep car'
9          /dev/sda1 20G 4.7G 15G 25% /etc/cars
10
11      student@lfs458-node-1a0a:~$ kubectl exec -it shell-
12      demo -- \
13          /bin/bash -c 'cat
14      /etc/cars/car.trim'
15      Shelby
```

11. Delete the Pod and ConfigMaps we were using.

```
1
2      student@lfs458-node-1a0a:~$ kubectl delete pods
3  shell-demo
4      pod "shell-demo" deleted
5
6      student@lfs458-node-1a0a:~$ kubectl delete configmap
7  fast-car colors
8      configmap "fast-car" deleted
9      configmap "colors" deleted
10
```

Exercise 9.2: Creating a Persistent NFS Volume (PV).

We will first deploy an NFS server. Once tested we will create a persistent NFS volume for containers to claim.

1. Install the software on your master node.

```
1
2      student@lfs458-node-1a0a:~$ sudo apt-get update &&
3  sudo \
4      apt-get install -y nfs-kernel-server
5      <output_omitted>
6
7
8
```

2. Make and populate a directory to be shared. Also give it similar permissions to **/tmp/**

```
1
2      student@lfs458-node-1a0a:~$ sudo mkdir /opt/sfw
3
4      student@lfs458-node-1a0a:~$ sudo chmod 1777
5  /opt/sfw/
6
7      student@lfs458-node-1a0a:~$ sudo bash -c \
8      'echo software > /opt/sfw/hello.txt'
9
10
```

3. Edit the NFS server file to share out the newly created directory. In this case we will share the directory with all. You can always **snoop** to see the inbound request in a later step and update the file to be more narrow.


```

1
2      student@lfs458-node-1a0a:~$ sudo vim /etc/exports
3      /opt/sfw/ *(rw, sync, no_root_squash, subtree_check)
4
5
6
7

```

4. Cause **/etc/exports** to be re-read:

```

1
2      student@lfs458-node-1a0a:~$ sudo exportfs -ra
3
4

```

5. Test by mounting the resource from your **second** node.

```

1
2      student@lfs458-worker:~$ sudo apt-get -y install
3  nfs-common
4      <output_omitted>
5
6      student@lfs458-worker:~$ showmount -e lfs458-node-
7  1a0a
8      Export list for lfs458-node-1a0a:
9      /opt/sfw *
10
11      student@lfs458-worker:~$ sudo mount
12  10.128.0.3:/opt/sfw /mnt
13
14      student@lfs458-worker:~$ ls -l /mnt
15      total 4
16      -rw-r--r-- 1 root root 9 Sep 28 17:55 hello.txt

```

6. Return to the master node and create a YAML file for the object with kind, PersistentVolume. Use the hostname of the master server and the directory you created in the previous step. Only syntax is checked, an incorrect name or directory will not generate an error, but a Pod using the resource will not start. Note that the accessModes do not currently affect actual access and are typically used as labels instead.

```

1
2      student@lfs458-node-1a0a:~$ vim PVol.yaml
3
4      apiVersion: v1
5      kind: PersistentVolume
6      metadata:
7          name: pvvol-1
8      spec:
9          capacity:
10             storage: 1Gi
11          accessModes:
12             - ReadWriteMany
13          persistentVolumeReclaimPolicy:
14      Retain
15
16          nfs:
17             path: /opt/sfw
18             server: lfs458-node-1a0a #<-- Edit to match
19      master node
20             readOnly: false
21

```

7. Create the persistent volume, then verify its creation.

```

1
2      student@lfs458-node-1a0a:~$ kubectl create -f
3      PVol.yaml
4      persistentvolume/pvvol-1 created
5
6      student@lfs458-node-1a0a:~$ kubectl get pv
7      NAME          CAPACITY          ACCESSMODES
8      RECLAIMPOLICY  STATUS
9      CLAIM          STORAGECLASS      REASON          AGE
10     pvvol-1          1Gi              RWX             Retain          Available
11
12     4s

```

Exercise 9.3: Creating a Persistent Volume Claim (PVC)

Before Pods can take advantage of the new PV we need to create a **Persistent Volume Claim (PVC)**.

1. Begin by determining if any currently exist.

```

1
2      student@lfs458-node-1a0a:~$ kubectl get pvc
3      No resources found.
4
5
6
7
8

```

2. Create a YAML file for the new pvc.

```

1
2      student@lfs458-node-1a0a:~$ vim pvc.yaml
3      apiVersion: v1
4      kind: PersistentVolumeClaim
5      metadata:
6          name: pvc-one
7      spec:
8      accessModes:
9          - ReadWriteMany
10         resources:
11             requests:
12                 storage: 200Mi
13
14
15

```

3. Create and verify the new pvc is bound. Note that the size is 1Gi, even though 200Mi was suggested. Only a volume of at least that size could be used.

```

1
2      student@lfs458-node-1a0a:~$ kubectl create -f
3      pvc.yaml
4      persistentvolumeclaim/pvc-one created
5
6      student@lfs458-node-1a0a:~$ kubectl get pvc
7      NAME      STATUS VOLUME CAPACITY ACCESSMODES
8      STORAGECLASS AGE
9      pvc-one Bound  pvvol-1 1Gi          RWX
10  4s
11

```

4. Look at the status of the pv again, to determine if it is in use. It should show a status of Bound.

```
1
2      student@lfs458-node-1a0a:~$ kubectl get pv
3
4      NAME      CAPACITY ACCESSMODES RECLAIMPOLICY STATUS
5  CLAIM                STORAGECLASS REASON AGE
6      pvvol-1 1Gi      RWX              Retain
7  Bound default/pvc-one
8  5m
9
10
```

5. Create a new deployment to use the pvc. We will copy and edit an existing deployment yaml file. We will change the deployment name then add a volumeMounts section under containers and volumes section to the general spec. The name used must match in both places, whatever name you use. The claimName must match an existing pvc. As shown in the following example.

```
1
2      student@lfs458-node-1a0a:~$ cp first.yaml nfs-
3 pod.yaml
4      student@lfs458-node-1a0a:~$ vim nfs-pod.yaml
5      apiVersion: apps/v1beta1
6      kind: Deployment
7      metadata:
8      annotations:
9      deployment.kubernetes.io/revision: "1"
10     generation: 1
11     labels:
12     run: nginx
13     name: nginx-nfs
14     namespace: default
15     resourceVersion: "1411"
16     spec:
17     replicas: 1
18     selector:
19     matchLabels:
20     run: nginx
21     strategy:
22     rollingUpdate:
23     maxSurge: 1
24     maxUnavailable: 1
25     type: RollingUpdate
26     template:
27     metadata:
28     creationTimestamp: null
29     labels:
30     run: nginx
31     spec:
32     containers:
33     - image: nginx
34     imagePullPolicy: Always
35     name: nginx
36     volumeMounts:
37     - name: nfs-vol
38     mountPath: /opt
39     ports:
40     - containerPort: 80
41     protocol: TCP
42     resources: {}
43     terminationMessagePath: /dev/termination-log
44     terminationMessagePolicy: File
45     volumes: #<<-- These four lines
46     - name: nfs-vol
47     persistentVolumeClaim:
48     claimName: pvc-one
49     dnsPolicy: ClusterFirst
```

```

50      restartPolicy: Always
51      schedulerName: default-scheduler
52      securityContext: {}
53      terminationGracePeriodSeconds: 30
54
55

```

6. Create the pod using the newly edited file

```

1
2      student@lfs458-node-1a0a:~$ kubectl create -f nfs-
3 pod.yaml
4      deployment.apps/nginx-nfs created
5
6

```

7. Look at the details of the pod. You may see the daemonset pods running as well.

```

1
2      student@lfs458-node-1a0a:~$ kubectl get pods
3      NAME                                     READY
4  STATUS RESTARTS AGE
5      nginx-nfs-1054709768-s8g28 1/1 Running 0
6  3m
7
8      student@lfs458-node-1a0a:~$ kubectl describe pod
9  nginx-nfs-1054709768-s8g28
10     Name:          nginx-nfs-1054709768-s8g28
11     Namespace:     default
12     Node:          lfs458-worker/10.128.0.5
13     <output_omitted>
14     Mounts:
15     /opt from nfs-vol (rw)
16     <output_omitted>
17     Volumes:
18         nfs-vol:
19             Type: PersistentVolumeClaim (a
20 reference to a PersistentV...
21             ClaimName: pvc-one
22             ReadOnly: false
23     <output_omitted>
24

```

8. View the status of the PVC. It should show as bound.

```

1
2      student@lfs458-node-1a0a:~$ kubectl get pvc
3      NAME      STATUS VOLUME CAPACITY ACCESS MODES
4  STORAGECLASS AGE
5      pvc-one Bound   pvvol-1 1Gi          RWX
6  2m
7

```

Exercise 9.4: Using a ResourceQuota to Limit PVC Count and Usage

The flexibility of cloud-based storage often requires limiting consumption among users. We will use the ResourceQuota object to both limit the total consumption as well as the number of persistent volume claims.

1. Begin by deleting the deployment we had created to use NFS, the pv and the pvc

```

1
2      student@lfs458-node-1a0a:~$ kubectl delete deploy
3  nginx-nfs
4      deployment.extensions "nginx-nfs" deleted
5
6      student@lfs458-node-1a0a:~$ kubectl delete pvc pvc-
7  one
8      persistentvolumeclaim "pvc-one" deleted
9
10     student@lfs458-node-1a0a:~$ kubectl delete pv pvvol-
11  1
12     persistentvolume "pvvol-1" deleted

```

2. Create a yaml file for the ResourceQuota object. Set the storage limit to ten claims with a total usage of 500Mi.

```

1
2      student@lfs458-node-1a0a:~$ vim storage-quota.yaml
3
4      apiVersion: v1
5      kind: ResourceQuota
6      metadata:
7          name: storagequota
8      spec:
9          hard:
10              persistentvolumeclaims: "10"
11              requests.storage: "500Mi"
12
13
14

```

3. Create a new namespace called small. View the namespace information prior to the new quota. Either the long name with double dashes --namespace or the nickname ns work for the resource.

```
1
2      student@lfs458-node-1a0a:~$ kubectl create namespace
3  small
4
5      namespace/small created
6      student@lfs458-node-1a0a:~$ kubectl describe ns
7  small
8      Name:                                small
9      Labels:                              <none>
10     Annotations:                          <none>
11     Status:                                Active
12
13     No resource quota.
14
15     No resource limits.
16
17
```

4. Create a new pv and pvc in the small namespace.

```
1
2      student@lfs458-node-1a0a:~$ kubectl create -f
3  PVol.yaml -n small
4      persistentvolume/pvvol-1 created
5
6      student@lfs458-node-1a0a:~$ kubectl create -f
7  pvc.yaml -n small
8      persistentvolumeclaim/pvc-one created
9
```

5. Create the new resource quota, placing this object into the low-usage-limit namespace.

```
1
2      student@lfs458-node-1a0a:~$ kubectl create -f
3  storage-quota.yaml \
4      -n small
5      resourcequota/storagequota created
6
7
8
```

6. Verify the small namespace has quotas. Compare the output to the same command above.


```
1
2      student@lfs458-node-1a0a:~$ kubectl describe ns
3  small
4      Name:                                small
5      Labels:
6      Annotations:
7      Status:                                Active
8
9      Resource Quotas
10     Name:                                storagequota
11     Resource                                Used Hard
12     -----                                --- ---
13
14     persistentvolumeclaims 1 10
15     requests.storage 200Mi 500Mi
16
17
18     No resource limits.
19
20
21
```

7. Remove the namespace line from the **nfs-pod.yaml** file. Should be around line 11 or so. This will allow us to pass other namespaces on the command line.

```
1
2      student@lfs458-node-1a0a:~$ vim nfs-pod.yaml
3
4
5
```

8. Create the container again.

```
1
2      student@lfs458-node-1a0a:~$ kubectl create -f nfs-
3  pod.yaml \
4      -n small
5      deployment.apps/nginx-nfs created
6
```

9. Determine if the deployment has a running pod.

```

1
2      student@lfs458-node-1a0a:~$ kubectl get deploy --
3 namespace=small
4      NAME                                DESIRED  CURRENT  UP-TO-DATE  AVAILABLE
5 AGE
6      nginx-nfs                           1        1        1           1
7 0      43s
8
9      student@lfs458-node-1a0a:~$ kubectl -n small
10 describe deploy \
11      nginx-nfs
12      <output_omitted>

```

10. Look to see if the pods are ready.

```

1
2      student@lfs458-node-1a0a:~$ kubectl get po -n small
3      NAME                                READY
4 STATUS RESTARTS AGE
5      nginx-nfs-2854978848-g3khf 1/1   Running 0
6 37s
7

```

11. . Ensure the Pod is running and is using the NFS mounted volume. If you pass the namespace first Tab will auto-complete the pod name.

```

1
2      student@lfs458-node-1a0a:~$ kubectl -n small
3 describe po \
4      nginx-nfs-2854978848-g3khf
5
6      Name:      nginx-nfs-2854978848-g3khf
7      Namespace: small
8      <output_omitted>
9
10      Mounts:
11      /opt from nfs-vol (rw)
12      <output_omitted>
13
14

```

12. View the quota usage of the namespace

```
1
2      student@lfs458-node-1a0a:~$ kubectl describe ns
3  small
4      <output_omitted>
5
6      Resource                               Quotas
7      Name:                                   storagequota
8      Resource      Used      Hard
9      -----      ---      ---
10
11     persistentvolumeclaims      1      10
12     requests.storage      200Mi
13     500Mi
14
15
16     No resource limits.
17
18
```

13. Create a 300M file inside of the /opt/sfw directory on the host and view the quota usage again. Note that with NFS the size of the share is not counted against the deployment.

```

1
2      student@lfs458-node-1a0a:~$ sudo dd if=/dev/zero \
3          of=/opt/sfw/bigfile bs=1M count=300
4      300+0 records in
5      300+0 records out
6      314572800 bytes (315 MB, 300 MiB) copied, 0.196794
7      s, 1.6 GB/s
8
9      student@lfs458-node-1a0a:~$ kubectl describe ns
10     small
11     <output_omitted>
12
13     Resource Quotas
14     Name:      storagequota
15     Resource   Used      Hard
16     -----   ---      ---
17
18     persistentvolumeclaims      1      10
19     requests.storage            200Mi
20     500Mi
21
22     <output_omitted>
23
24     student@lfs458-node-1a0a:~$ du -h /opt/
25     301M    /opt/sfw
26     41M     /opt/cni/bin
27     41M     /opt/cni
28     341M    /opt/
29

```

14. Now let us illustrate what happens when a deployment requests more than the quota. Begin by shutting down the existing deployment.

```

1
2      student@lfs458-node-1a0a:~$ kubectl -n small get
3     deploy
4     NAME          DESIRED CURRENT UP-TO-DATE AVAILABLE
5     AGE
6     nginx-nfs      1          1          1
7     1             11m
8
9      student@lfs458-node-1a0a:~$ kubectl -n small delete
10     deploy nginx-nfs
11     deployment.extensions "nginx-nfs" deleted

```

15. Once the Pod has shut down view the resource usage of the namespace again. Note the storage did not get cleaned up when the pod was shut down.

```

1
2      student@lfs458-node-1a0a:~$ kubectl describe ns
3  small
4      <output_omitted>
5
6      Resource                               Quotas
7      Name:                                   storagequota
8      Resource      Used      Hard
9      -----      ---      ---
10
11     persistentvolumeclaims      1      10
12     requests.storage      200Mi
13     500Mi

```

16. Remove the pvc then view the pv it was using. Note the RECLAIM POLICY and STATUS

```

1
2      student@lfs458-node-1a0a:~$ kubectl get pvc -n small
3      NAME      STATUS  VOLUME  CAPACITY  ACCESSMODES
4  STORAGECLASS AGE
5      pvc-one Bound   pvvol-1 1Gi      RWX
6  19m
7
8      student@lfs458-node-1a0a:~$ kubectl -n small delete
9  pvc pvc-one
10     persistentvolumeclaim "pvc-one" deleted
11
12     student@lfs458-node-1a0a:~$ kubectl -n small get pv
13     NAME      CAPACITY  ACCESSMODES
14  RECLAIMPOLICY STATUS  CLAIM
15  STORAGECLASS REASON AGE
16     pvvol-1      1Gi RWX Retain Released small/pvc-one
17  44m

```

17. Dynamically provisioned storage uses the ReclaimPolicy of the StorageClass which could be Delete, Retain, or some types allow Recycle. Manually created persistent volumes default to Retain unless set otherwise at creation. The default storage policy is to retain the storage to allow recovery of any data. To change this begin by viewing the yaml output.

```

1
2      student@lfs458-node-1a0a:~$ kubectl get pv/pvvol-1 -
3  o yaml
4      ....
5                  path: /opt/sfw
6                  server: lfs458-node-1a0a
7                  persistentVolumeReclaimPolicy: Retain
8      status:
9      phase: Released
10
11
12

```

18. Currently we will need to delete and re-create the object. Future development on a deleter plugin is planned. We will re-create the volume and allow it to use the Retain policy, then change it once running.

```

1
2      student@lfs458-node-1a0a:~$ kubectl delete pv/pvvol-
3  1
4      persistentvolume "pvvol-1" deleted
5
6      student@lfs458-node-1a0a:~$ grep Retain PVol.yaml
7      persistentVolumeReclaimPolicy: Retain
8
9      student@lfs458-node-1a0a:~$ kubectl create -f
10 PVol.yaml
11      persistentvolume "pvvol-1" created
12
13

```

19. We will use kubectl patch to change the retention policy to Delete. The yaml output from before can be helpful in getting the correct syntax

```

1
2      student@lfs458-node-1a0a:~$ kubectl patch pv pvvol-1
3  -p \
4      '{"spec":
5      {"persistentVolumeReclaimPolicy":"Delete"}}'
6
7      persistentvolume/pvvol-1 patched
8      student@lfs458-node-1a0a:~$ kubectl get pv/pvvol-1
9      NAME          CAPACITY          ACCESSMODES
10 RECLAIMPOLICY     STATUS    CLAIM
11                STORAGECLASS REASON AGE
12      pvvol-1 1Gi RWX Delete Available 2m

```

20. View the current quota settings

```

1
2      student@lfs458-node-1a0a:~$ kubectl describe ns small
3      .
4      requests.storage          0          500Mi
5
6
7

```

21. Create the pvc again. Even with no pods running, note the resource usage.

```

1
2      student@lfs458-node-1a0a:~$ kubectl -n small create
3      -f pvc.yaml
4      persistentvolumeclaim/pvc-one created
5
6      student@lfs458-node-1a0a:~$ kubectl describe ns
7      small
8      .
9      requests.storage          200Mi
10     500Mi
11

```

22. Remove the existing quota from the namespace.

```

1
2      student@lfs458-node-1a0a:~$ kubectl -n small get
3      resourcequota
4      NAME          CREATED AT
5      storagequota  2018-08-01T04:10:02Z
6
7      student@lfs458-node-1a0a:~$ kubectl -n small delete
8      \
9      resourcequota storagequota
10     resourcequota "storagequota" deleted
11
12

```

23. Edit the storagequota.yaml file and lower the capacity to 100Mi.

```

1
2      student@lfs458-node-1a0a:~$ vim storage-quota.yaml
3      .
4      requests.storage: "100Mi"
5
6
7

```

24. Create and verify the new storage quota. Note the hard limit has already been exceeded

```

1
2      student@lfs458-node-1a0a:~$ kubectl create -f
3 storage-quota.yaml -n small
4      resourcequota/storagequota created
5
6      student@lfs458-node-1a0a:~$ kubectl describe ns
7 small
8      .
9      persistentvolumeclaims          1
10 10
11      requests.storage                200Mi
100Mi
      No resource limits.
```

25. Create the deployment again. View the deployment. Note there are no errors seen.

```

1
2      student@lfs458-node-1a0a:~$ kubectl create -f nfs-
3 pod.yaml \
4      -n small
5      deployment.apps/nginx-nfs created
6
7      student@lfs458-node-1a0a:~$ kubectl -n small
8 describe deploy/nginx-nfs
9      Name:          nginx-nfs
10     Namespace:    small
11     <output_omitted>
12
```

26. Examine the pods to see if they are actually running.

```

1
2      student@lfs458-node-1a0a:~$ kubectl -n small get po
3      NAME                                                    READY
4  STATUS RESTARTS AGE
5      nginx-nfs-2854978848-vb6bh  1/1    Running 0      58s
6
7
```

27. As we were able to deploy more pods even with apparent hard quota set, let us test to see if the reclaim of storage takes place. Remove the deployment and the persistent volume claim.


```

1
2      student@lfs458-node-1a0a:~$ kubectl -n small delete
3  deploy nginx-nfs
4      deployment.extensions "nginx-nfs" deleted
5
6      student@lfs458-node-1a0a:~$ kubectl -n small delete
7  pvc/pvc-one
8      persistentvolumeclaim "pvc-one" deleted
9

```

28. View if the persistent volume exists. You will see it attempted a removal, but failed. If you look closer you will find the error has to do with the lack of a deleter volume plugin for NFS. Other storage protocols have a plugin.

```

1
2      student@lfs458-node-1a0a:~$ kubectl -n small get pv
3      NAME                                CAPACITY
4  ACCESSMODES                            RECLAIMPOLICY          STATUS
5  CLAIM
6
7      REASON                                STORAGECLASS           AGE
8      pvvol-1 1Gi RWX Delete Failed small/pvc-one 20m

```

29. Ensure the deployment, pvc and pv are all removed.

```

1
2      student@lfs458-node-1a0a:~$ kubectl delete pv/pvvol-1
3      persistentvolume "pvvol-1" deleted
4
5
6

```

30. Edit the persistent volume YAML file and change the persistentVolumeReclaimPolicy: to Recycle.

```

1
2      student@lfs458-node-1a0a:~$ vim PVol.yaml
3      ....
4      persistentVolumeReclaimPolicy: Recycle
5      ....
6
7
8
9

```

31. Add a LimitRange to the namespace and attempt to create the persistent volume and persistent volume claim again. We can use the **LimitRange** we used earlier

```

1
2      student@lfs458-node-1a0a:~$ kubectl -n small create -
3  f \
4      low-resource-range.yaml
5      limitrange/low-resource-range created
6
7

```

32. View the settings for the namespace. Both quotas and resource limits should be seen.

```

1
2      student@lfs458-node-1a0a:~$ kubectl describe ns
3  small
4      <output_omitted>
5      Resource Limits
6
7      Type      Resource      Min      Max      Default
8  Request      Default Limit  ...
9  ---          -...
10 Container cpu - - 500m 1 -
11 Container memory - - 100Mi 500Mi -

```

33. Create the persistent volume again. View the resource. Note the Reclaim Policy is Recycle.

```

1
2      student@lfs458-node-1a0a:~$ kubectl -n small create
3  -f PVol.yaml
4      persistentvolume/pvvol-1 created
5
6      student@lfs458-node-1a0a:~$ kubectl get pv
7      NAME      CAPACITY ACCESS MODES RECLAIM POLICY STATUS
8  ...
9      pvvol-1 1Gi      RWX      Recycle
10 Available ...

```

34. Attempt to create the persistent volume claim again. The quota only takes effect if there is also a resource limit in effect.

```
1
2      student@lfs458-node-1a0a:~$ kubectl -n small create
3 -f pvc.yaml
4      Error from server (Forbidden): error when creating
5 "pvc.yaml":
6      persistentvolumeclaims "pvc-one" is forbidden:
7 exceeded quota:
8      storagequota, requested: requests.storage=200Mi,
9      used:
10     requests.storage=0, limited: requests.storage=100Mi
```

35. Edit the resourcequota to increase the requests.storage to 500mi

```
1
2      student@lfs458-node-1a0a:~$ kubectl -n small edit
3 resourcequota
4      ....
5          spec:
6              hard:
7                  persistentvolumeclaims: "10"
8          requests.storage: 500Mi
9          status:
10             hard:
11                 persistentvolumeclaims: "10"
12             ....
13
14
15
```

36. Create the pvc again. It should work this time. Then create the deployment again

```
1
2      student@lfs458-node-1a0a:~$ kubectl -n small create -
3 f pvc.yaml
4      persistentvolumeclaim/pvc-one created
5
6      student@lfs458-node-1a0a:~$ kubectl -n small create -
7 f nfs-pod.yaml
8      deployment.apps/nginx-nfs created
9
```

37. View the namespace settings.

```

1
2      student@lfs458-node-1a0a:~$ kubectl describe ns small
3      <output_omitted>
4
5
6
7

```

38. Delete the deployment. View the status of the **pv** and **pvc**

```

1
2      student@lfs458-node-1a0a:~$ kubectl -n small delete
3      deploy nginx-nfs
4      deployment.extensions "nginx-nfs" deleted
5
6      student@lfs458-node-1a0a:~$ kubectl get pvc -n small
7      NAME      STATUS VOLUME CAPACITY ACCESS MODES
8      STORAGECLASS AGE
9      pvc-one Bound  pvvol-1 1Gi          RWX
10     7m
11
12     student@lfs458-node-1a0a:~$ kubectl -n small get pv
13     NAME CAPACITY ACCESS MODES RECLAIM POLICY STATUS
14     CLAIM STORA...
15     pvvol-1 1Gi RWX Recycle Bound small/pvc-one

```

39. Delete the **pvc** and check the status of the pv. It should show as Available

```

1
2      student@lfs458-node-1a0a:~$ kubectl -n small delete
3      pvc pvc-one
4      persistentvolumeclaim "pvc-one" deleted
5
6      student@lfs458-node-1a0a:~$ kubectl -n small get pv
7      NAME      CAPACITY ACCESS MODES RECLAIM POLICY STATUS
8      CLAIM STORA...
9      pvvol-1 1Gi          RWX          Recycle
10     Available ...

```

40. Remove the pv and any other resources created during this lab.

```
1
2      student@lfs458-node-1a0a:~$ kubectl delete pv pvvol-1
3      persistentvolume "pvvol-1" deleted
4
5
6
7
```

Reference of this article can be found below:-

Kubernetes Courses | Training | Trainer | Job Support | Kubernetes Consultant

Avail Rajesh Kumar as trainer at

50% Discount

Puppet Online Training

Enroll Now

Puppet Classroom Training

Enroll Now

What do you think?

0 Responses



DevOpsSchool Comment Policy

<http://www.devopsschool.com/>

Please read our [Comment Policy](#) before commenting.



0 Comments

DevOpsSchool

Disqus' Privacy Policy

Login ▾

Recommend

Tweet

Share

Sort by Best ▾



Start the discussion...

LOG IN WITH

OR SIGN UP WITH DISQUS

Name

Be the first to comment.

Subscribe Add Disqus to your siteAdd DisqusAdd Do Not Sell My Data

